

# Quantitative Economics Workshop Paris

## Prelude to Dynamic Programming

John Stachurski

September 2022

# Introduction

Summary of this lecture:

- Linear equations
- Fixed point theory
- Algorithms
- Job search

Operations on real numbers such as  $|\cdot|$  and  $\vee$  are applied to vectors element-by-element

Example.

$$a = \begin{pmatrix} a_1 \\ \vdots \\ a_n \end{pmatrix} \implies |a| = \begin{pmatrix} |a_1| \\ \vdots \\ |a_n| \end{pmatrix}$$

$$a \vee b = \begin{pmatrix} a_1 \vee b_1 \\ \vdots \\ a_n \vee b_n \end{pmatrix} \quad \text{and} \quad a \wedge b = \begin{pmatrix} a_1 \wedge b_1 \\ \vdots \\ a_n \wedge b_n \end{pmatrix}$$

etc.

# Linear Equations

Given one-dimensional equation  $x = ax + b$ , we have

$$|a| < 1 \quad \implies \quad x^* = \frac{b}{1-a} = \sum_{k \geq 0} a^k b$$

How can we extend this beyond one dimension?

- When does  $x = Ax + b$  have a unique solution?
- How can we compute it?

Recall that  $\lambda \in \mathbb{C}$  is an eigenvalue of  $n \times n$  matrix  $A$  if

$$Ae = \lambda e$$

for some nonzero  $e \in \mathbb{C}^n$ .

We define the **spectral radius** of square matrix  $A$  as

$$r(A) := \max\{|\lambda| : \lambda \text{ is an eigenvalue of } A\}$$

**Key idea:**

- $r(A) < 1$  is a generalization of  $|a| < 1$

# Neumann Series Lemma

Suppose

- $b$  is  $n \times 1$  and  $A$  is  $n \times n$
- $I$  is the  $n \times n$  identity matrix

**Theorem.** If  $r(A) < 1$ , then

1.  $I - A$  is nonsingular
2.  $\sum_{k \geq 0} A^k$  converges
3.  $(I - A)^{-1} = \sum_{k \geq 0} A^k$  and
4.  $x = Ax + b$  has the unique solution

$$x^* := (I - A)^{-1}b$$

Intuitive idea: with  $S := \sum_{k \geq 0} A^k$ , we have

$$\begin{aligned} I + AS &= I + A(I + A + \dots) \\ &= I + A + A^2 + \dots \\ &= S \end{aligned}$$

Rearranging  $I + AS = S$  gives  $S = (I - A)^{-1}$

$$\therefore x = Ax + b \iff (I - A)x = b \iff x^* = (I - A)^{-1}b$$

# Fixed Points

To solve nonlinear equations we use fixed point theory

Recall: if  $S$  is any set then

- $T$  is a **self-map** on  $S$  if  $T$  maps  $S$  into itself
- $x^* \in S$  is called a **fixed point** of  $T$  in  $S$  if  $Tx^* = x^*$

Examples.

- Every  $x$  in  $S$  is fixed for the **identity map**  $I: x \mapsto x$
- If  $S = \mathbb{N}$  and  $Tx = x + 1$ , then  $T$  has no fixed point



## Examples.

- If  $S = \mathbb{R}$  and  $Tx = x^2$ , then  $T$  has fixed points at  $0, 1$
- If  $S \subset \mathbb{R}$ ,  $Tx = x \iff T$  meets the 45 degree line

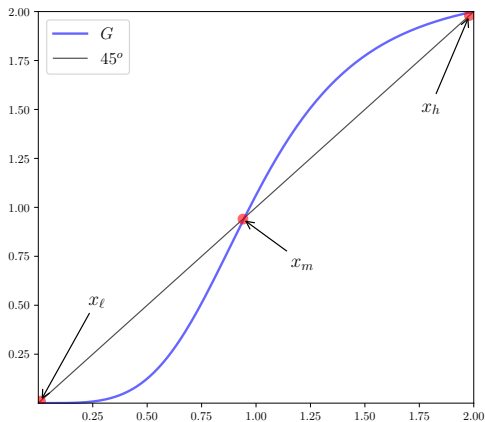


Figure: Graph and fixed points of  $G: x \mapsto 2.125/(1+x^{-4})$

**Key idea:** Fixed point theory is for solving equations

1. solve  $x = Tx \iff$  find fixed points of  $T$
2. solve  $Gx = 0 \iff$  find fixed points of  $Tx := Gx + x$

**Example.** If  $S = \mathbb{R}^n$  and  $Tx = Ax + b$ , then

$x^*$  solves equation  $x = Ax + b \iff x^*$  is a fixed point of  $T$

(But fixed point theory is mainly for nonlinear equations)

Point on notation:

- $T^2 = T \circ T$
- $T^3 = T \circ T \circ T$
- etc.

**Example.**  $Tx = Ax + b$  implies  $T^2x = A(Ax + b) + b$

Self-map  $T$  is called **globally stable** on  $S$  if

1.  $T$  has a unique fixed point  $x^*$  in  $S$  and
2.  $T^k x \rightarrow x^*$  as  $k \rightarrow \infty$  for all  $x \in S$

**Example.** Let  $S = \mathbb{R}^n$  and  $Tx = Ax + b$

**Ex.** Prove:  $r(A) < 1$  implies  $T$  is globally stable on  $S$

**Example.** Consider Solow–Swan growth dynamics

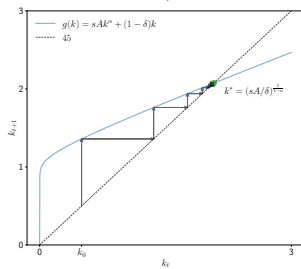
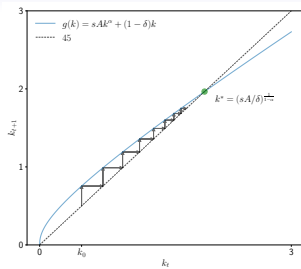
$$k_{t+1} = g(k_t) := sAk_t^\alpha + (1 - \delta)k_t, \quad t = 0, 1, \dots,$$

where

- $k_t$  is capital stock per worker,
- $A, \alpha > 0$  are production parameters,  $\alpha < 1$
- $s > 0$  is a savings rate, and
- $\delta \in (0, 1)$  is a rate of depreciation

Iterating with  $g$  from  $k_0$  generates a time path for capital stock

The map  $g$  is globally stable on  $(0, \infty)$



## Note from last slide

- If  $g$  is flat near  $k^*$ , then  $g(k) \approx k^*$  for  $k$  near  $k^*$
- A flat function near the fixed point  $\implies$  fast convergence

## Conversely

- If  $g$  is close to the 45 degree line near  $k^*$ , then  $g(k) \approx k$
- Close to 45 degree line means high persistence, slow convergence



Given a self-map  $T$  on  $S$ , we typically ask

- Does  $T$  have at least one fixed point on  $S$  (existence)?
- Does  $T$  have at most one fixed point on  $S$  (uniqueness)?
- How can we compute fixed points of  $T$ ?

For the last question, we seek an algorithm

Then we investigate its properties

# Contractions

Let

- $U$  be a nonempty subset of  $\mathbb{R}^n$ ,
- $\|\cdot\|$  be a norm on  $\mathbb{R}^n$ , and
- $T$  be a self-map on  $U$

$T$  is called a **contraction** on  $U$  with respect to  $\|\cdot\|$  if

$$\exists \lambda < 1 \text{ such that } \|Tu - Tv\| \leq \lambda \|u - v\| \quad \text{for all } u, v \in U$$

**Example.**  $Tx = ax + b$  is a contraction on  $\mathbb{R}$  with respect to  $|\cdot|$  if and only if  $|a| < 1$

Indeed,

$$|Tx - Ty| = |ax + b - ay - b| = |a||x - y|$$

# Banach's Contraction Mapping Theorem

## Theorem If

1.  $U$  is closed in  $\mathbb{R}^n$  and
2.  $T$  is a contraction of modulus  $\lambda$  on  $U$  with respect to some norm  $\|\cdot\|$  on  $\mathbb{R}^n$ ,

then  $T$  has a unique fixed point  $u^*$  in  $U$  and

$$\|T^n u - u^*\| \leq \lambda^n \|u - u^*\| \quad \text{for all } n \in \mathbb{N} \text{ and } u \in U$$

In particular,  $T$  is globally stable on  $U$

Proof: See the DP text

# Successive approximation

---

fix a guess  $x_0$  and some error tolerance  $\tau$

$\varepsilon \leftarrow \tau + 1$

$x \leftarrow x_0$

**while**  $\varepsilon > \tau$  **do**

$y \leftarrow Tx$

$\varepsilon \leftarrow \|y - x\|$

$x \leftarrow y$

**end**

**return**  $x$

---

If  $T$  is a contraction, say, then the output will be close to  $x^*$

---

```
import numpy as np
def successive_approx(T,                                # Operator (callable)
                      x_0,                             # Initial condition
                      tolerance=1e-6,                  # Error tolerance
                      max_iter=10_000,                 # Max iteration bound
                      print_step=25,                  # Print at multiples
                      verbose=False):

    x = x_0
    error = tolerance + 1
    k = 1
    while error > tolerance and k <= max_iter:
        x_new = T(x)
        error = np.max(np.abs(x_new - x))
        if verbose and k % print_step == 0:
            print(f"Completed iteration {k} with error {error}.")
        x = x_new
        k += 1
    if error > tolerance:
        print(f"Warning: Iteration hit upper bound {max_iter}.")
    elif verbose:
        print(f"Terminated successfully in {k} iterations.")
    return x
```




Figure: Successive approximation from different initial conditions

# Newton's Method

Let  $g$  be a smooth self-map on  $S := (a, b)$

We start with guess  $x_0$  of the fixed point and update it using

$$g_a(x) \approx g(x_0) + g'(x_0)(x - x_0)$$

Set  $g_a(x_1) = x_1$  and solve for  $x_1$  to get

$$x_1 = \frac{x_0 - g'(x_0)x_0}{1 - g'(x_0)}$$

Now repeat logic to get  $x_2, x_3, \dots$

This gives **Newton's method**

We are iterating on the map

$$x_{k+1} = q(x_k) \quad \text{where} \quad q(x) := \frac{x - g'(x)x}{1 - g'(x)}$$

Note that we are applying successive approximation to  $q$

Hence we can use the same code



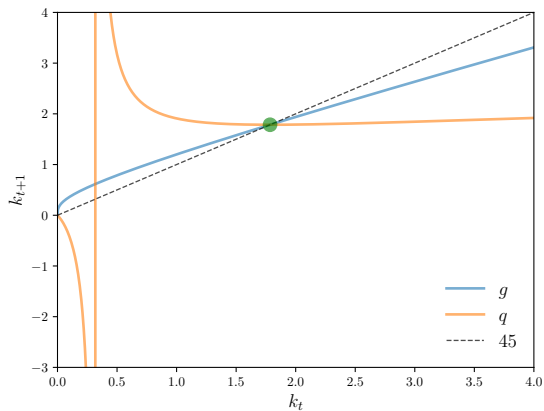


Figure: Successive approximation vs Newton's method

## Comments:

- The map  $q$  is flat close to the fixed point  $k^*$
- Hence Newton's method converges quickly near  $k^*$
- But Newton's method is not globally convergent
- Successive approximation is slower but more robust

## Key ideas

- There is almost always a trade-off between robustness and speed
- Speed requires assumptions, and assumptions can fail

**Newton's method** extends naturally to **multiple dimensions**

When  $h$  is a map from  $S \subset \mathbb{R}^n$  to itself, we use

$$x_{k+1} = x_k - [J(x_k)]^{-1}h(x_k)$$

Here  $J_h(x_k) :=$  the Jacobian of  $h$  evaluated at  $x_k$

Comments

- Typically faster but less robust
- Matrix operations can be parallelized
- Automatic differentiation can be helpful

# Job Search

A model of job search created by **John J. McCall**

We model the decision problem of an unemployed worker

Job search depends on

- current and likely future wage offers
- impatience, and
- the availability of unemployment compensation

We begin with a very simple version of the McCall model

(Later we consider extensions)

# Set Up

An agent begins working life at time  $t = 0$  without employment

Receives a new job offer paying wage  $w_t$  at each date  $t$

She has two choices:

1. **accept** the offer and work permanently at  $w_t$  or
2. **reject** the offer, receive unemployment compensation  $c$ , and reconsider next period

Assume  $\{w_t\}$  is  $\overset{\text{IID}}{\sim} \varphi$ , where

- $W \subset \mathbb{R}_+$  is a finite set of wage outcomes and
- $\varphi \in \mathcal{D}(W)$

The agent cares about the future but is **impatient**

Impatience is parameterized by a **time discount factor**  $\beta \in (0,1)$

- Present value of a next-period payoff of  $y$  dollars is  $\beta y$

Trade off:

- $\beta < 1$  indicating some impatience
- hence the agent will be tempted to accept reasonable offers, rather than always waiting for a better one
- The key question is how long to wait

The worker who aims to maximize

$$\mathbb{E} \sum_{t=0}^{\infty} \beta^t Y_t, \quad Y_t \in \{c, W_t\} \text{ is earnings at time } t \quad (1)$$

- $\{W_t\} \stackrel{\text{iid}}{\sim} \varphi$  for  $\varphi \in \mathcal{D}(W)$
- $W \subset \mathbb{R}_+$  with  $|W| < \infty$
- $c$  and  $\beta$  are positive and  $\beta < 1$
- jobs are permanent

What is max EPV of each option when lifetime is infinite?

What if we **accept**  $w \in W$  now?

$$\text{EPV} = \text{stopping value} = w + \beta w + \beta^2 w + \dots = \frac{w}{1 - \beta}$$

What if we **reject**?

EPV = continuation value

= EPV of optimal choice in each subsequent period

But what are optimal choices?!

Calculating optimal choice requires knowing optimal choice!



# The Value Function

Let  $v^*(w) := \max$  lifetime EPV given wage offer  $w$

We call  $v^*$  the **value function**

**Suppose** that we know  $v^*$

Then the (maximum) **continuation value** is

$$h^* := c + \beta \sum_{w' \in W} v^*(w') \varphi(w')$$

= max EPV conditional on decision to continue

The optimal choice is then

$$\mathbb{1} \{ \text{stopping value} \geq \text{continuation value} \} = \mathbb{1} \left\{ \frac{w}{1 - \beta}, h^* \right\}$$

But how can we calculate  $v^*$ ?

**Key idea:** We can use the Bellman equation to solve for  $v^*$

**Theorem.** The value function  $v^*$  satisfies the **Bellman equation**

$$v^*(w) = \max \left\{ \frac{w}{1 - \beta}, c + \beta \sum_{w' \in W} v^*(w') \phi(w') \right\} \quad (w \in W)$$

Intuition:

- If accept, get  $w/(1 - \beta)$
- If reject and then choose optimally, get max continuation value
- Max value today is max of these alternatives

So how can we use the Bellman equation

$$v^*(w) = \max \left\{ \frac{w}{1 - \beta}, c + \beta \sum_{w' \in W} v^*(w') \varphi(w') \right\} \quad (w \in W)$$

to solve for  $v^*$ ?

Let's now return to the job search problem

Recall that the value function  $v^*$  solves the Bellman equation

That is,

$$v^*(w) = \max \left\{ \frac{w}{1 - \beta}, c + \beta \sum_{w' \in W} v^*(w') \varphi(w') \right\} \quad (w \in W)$$

The infinite-horizon **continuation value** is defined as

$$h^* := c + \beta \sum_{w'} v^*(w') \varphi(w')$$

**Key question:** how to solve for  $v^*$ ?

We introduce the **Bellman operator**, defined at  $v \in \mathbb{R}^W$  by

$$(Tv)(w) = \max \left\{ \frac{w}{1-\beta}, c + \beta \sum_{w' \in W} v(w') \varphi(w') \right\} \quad (w \in W)$$

By construction,  $Tv = v \iff v$  solves the Bellman equation

Let  $\mathcal{V} := \mathbb{R}_+^W$

**Proposition.**  $T$  is a contraction  $\mathcal{V}$  with respect to  $\|\cdot\|_\infty$

In the proof, we use the elementary bound

$$|\alpha \vee x - \alpha \vee y| \leq |x - y| \quad (\alpha, x, y \in \mathbb{R})$$

Fixing  $f, g$  in  $\mathcal{V}$  fix any  $w \in W$ , we have

$$\begin{aligned} |(Tf)(w) - (Tg)(w)| &\leq \left| \beta \sum_{w'} f(w') \varphi(w') - \beta \sum_{w'} g(w') \varphi(w') \right| \\ &= \beta \left| \sum_{w'} [f(w') - g(w')] \varphi(w') \right| \end{aligned}$$

Applying the triangle inequality,

$$|(Tf)(w) - (Tg)(w)| \leq \beta \sum_{w'} |f(w') - g(w')| \varphi(w') \leq \beta \|f - g\|_{\infty}$$

$$\therefore \|Tf - Tg\|_{\infty} \leq \beta \|f - g\|_{\infty}$$

Recall: The optimal decision at any given time, facing current wage draw  $w \in W$ , is

$$\mathbb{1} \left\{ \frac{w}{1 - \beta} \geq h^* \right\}$$

Let's try to write this in the language of dynamic programming

Dynamic programming centers around the problem of finding optimal **policies**

# Optimal Policies

In general, for a dynamic program, choices consist of a sequence  $(A_t)_{t \geq 0}$

- specifies how the agent acts at each  $t$

Since agents are not clairvoyant, so we assume that  $A_t$  cannot depend on future events

In other words, for some function  $\sigma_t$ ,

$$A_t = \sigma_t(X_t, A_{t-1}, X_{t-1}, A_{t-2}, X_{t-2}, \dots, A_0, X_0)$$

In dynamic programming,  $\sigma_t$  is called a **policy function**



**Key idea** Design the state such that  $X_t$  is

- sufficient to determine the optimal current action
- but not so large as to be unmanagable
- Finding the state is an art!

**Example.** Recall retailer who chooses stock orders and prices in each period

What to include in the current state?

- level of current inventories
- interest rates and inflation?
- the rate at which inventories have changed?
- competitors prices?

So suppose state  $X_t$  determines the current action  $A_t$

Then we can write  $A_t = \sigma(X_t)$  for some function  $\sigma$

Note that we dropped the time subscript on  $\sigma$

No loss of generality: can include time in the current state

- i.e., expand  $X_t$  to  $\hat{X}_t = (t, X_t)$

Depends on the problem at hand

- For the job search model with finite horizon, the date matters
- For the infinite horizon version of the problem, however, the agent always looks forward toward an infinite horizon

For job search model,

- state = current wage offer and
- possible actions are accept (1) or reject (0)

A policy is a map  $\sigma$  from  $W$  to  $\{0, 1\}$

Let  $\Sigma$  be the set of all such maps

For each  $v \in \mathcal{V}$ , let us define a  **$v$ -greedy policy** to be a  $\sigma \in \Sigma$  satisfying

$$\sigma(w) = \mathbb{1} \left\{ \frac{w}{1 - \beta} \geq c + \beta \sum_{w' \in W} v(w') \varphi(w') \right\} \quad \text{for all } w \in W$$

Accepts iff  $w/(1 - \beta) \geq$  continuation value computed using  $v$

Optimal choice:

- agent should adopt a  $v^*$ -greedy policy
- Sometimes called **Bellman's principle of optimality**

We can also express a  $v^*$ -greedy policy via

$$\sigma^*(w) = \mathbb{1} \{w \geq w^*\} \quad \text{where } w^* := (1 - \beta)h^* \quad (2)$$

The term  $w^*$  in (2) is called the **reservation wage**

- Same ideas as before, different language
- We prove optimality more carefully later

# Computation

Since  $T$  is globally stable on  $\mathcal{V}$ , we can compute an approximate optimal policy by

1. applying successive approximation on  $T$  to compute  $v^*$
2. calculate a  $v^*$ -greedy policy

In dynamic programming, this approach is called **value function iteration**

---

---

input  $v_0 \in \mathcal{V}$ , an initial guess of  $v^*$

input  $\tau$ , a tolerance level for error

$\varepsilon \leftarrow \tau + 1$

$k \leftarrow 0$

**while**  $\varepsilon > \tau$  **do**

**for**  $w \in W$  **do**

$v_{k+1}(w) \leftarrow (Tv_k)(w)$

**end**

$\varepsilon \leftarrow \|v_k - v_{k+1}\|_\infty$

$k \leftarrow k + 1$

**end**

Compute a  $v_k$ -greedy policy  $\sigma$

**return**  $\sigma$

---





Figure: A sequence of iterates of the Bellman operator



**Figure:** The approximate value function for job search



# Computing the Continuation Value Directly

We used a standard dynamic programming approach to solve this problem

Sometimes we can find more efficient ways to solve particular problems

For the infinite horizon job search problem, a more efficient way exists

The idea is to compute the continuation value directly

This shifts the problem from  $n$ -dimensional to one-dimensional

Method: Recall that

$$v^*(w) = \max \left\{ \frac{w}{1 - \beta}, c + \beta \sum_{w'} v^*(w') \varphi(w') \right\} \quad (w \in W)$$

Using the definition of  $h^*$ , we can write

$$v^*(w') = \max \{ w' / (1 - \beta), h^* \} \quad (w' \in W)$$

Take expectations, multiply by  $\beta$  and add  $c$  to obtain

$$h^* = c + \beta \sum_{w'} \max \left\{ \frac{w'}{1 - \beta}, h^* \right\} \varphi(w')$$

How to find  $h^*$  from the equation

$$h^* = c + \beta \sum_{w'} \max \left\{ \frac{w'}{1 - \beta}, h^* \right\} \varphi(w') \quad (3)$$

We introduce the map  $g: \mathbb{R}_+ \rightarrow \mathbb{R}_+$  defined by

$$g(h) = c + \beta \sum_{w'} \max \left\{ \frac{w'}{1 - \beta}, h \right\} \varphi(w')$$

By construction,  $h^*$  solves (3) if and only if  $h^*$  is a fixed point of  $g$

**Ex.** Show that  $g$  is a contraction map on  $\mathbb{R}_+$




Figure: Computing the continuation value as the fixed point of  $g$

New algorithm:

1. Compute  $h^*$  via successive approximation on  $g$ 
  - Iteration in  $\mathbb{R}$ , not  $\mathbb{R}^n$
2. Optimal policy is

$$\sigma^*(w) = \mathbb{1} \left\{ \frac{w}{1-\beta} \geq h^* \right\}$$